

Unfolding the Leech Lattice: Fused Multi-Shell Decoding and VRAM Layouts for 2-Bit LLM Weights

What a combinatorial index costs to serve, measured against deployed 4-bit and 2-bit GEMV kernels

Author

Pier-Jean Malandrino
Scub, Bordeaux, France
pierjean.malandrino@scub.net

Corresponding author, and sole author.

Manuscript type

New Paper, Not an Extension of a Conference Paper. The work has not been published before in any form and is not under review elsewhere.

Acknowledgements

The measurement campaign was self-funded; the GPU time was rented per job and its cost is itemized in the artifact repository. The author thanks the maintainers of the open-source projects this work builds on and measures against.

Artifact availability

The manuscript's availability section withholds these URLs for double-anonymous review. They are, for the editor:

Code, data and paper source: <https://github.com/pjmalandrino/llvq>

Quantized 4B artifact: <https://huggingface.co/Pier-Jean/Qwen3-4B-LLVQ-2bit>

The repository is set private for the duration of the review, so that a search on the manuscript's distinctive strings cannot identify the author. Access can be granted to the editor or to reviewers on request.

Use of generative AI

Disclosed under ACM's Policy on Authorship, and repeated in the manuscript itself. Anthropic's Claude, through the Claude Code CLI in 2026, drafted and revised the manuscript's text, wrote the figure-generating and table-checking scripts, wrote portions of the CUDA kernels, Rust host code and tests, and recomputed statistics from the committed measurement journals. The author designed the study, authorized every measurement, verified each number against its source, and is responsible for the entire content.