

Unfolding the Leech Lattice: Fused Multi-Shell Decoding and VRAM Layouts for 2-Bit LLM Weights

What a combinatorial index costs to serve, measured against deployed 4-bit and 2-bit GEMV kernels

Author

Pier-Jean Malandrino
Scub, Dijon, France
pierjean.malandrino@scub.net

Corresponding author, and sole author.

Manuscript type

New Paper, Not an Extension of a Conference Paper. The work has not been published before in any form and is not under review elsewhere.

Acknowledgements

The measurement campaign was self-funded; the GPU time was rented per job and its cost is itemized in the artifact repository. The author thanks the maintainers of the open-source projects this work builds on and measures against.

Artifact availability

The manuscript's availability section withholds the repository and model URLs for double-anonymous review. They can be supplied to the editor on request.